

I/O in Linux

I/O in Linux: Kernel I/O Structure

What Is I/O?

Input/output (I/O) refers to the process of transferring data between a computer's internal memory [central processing unit (CPU), RAM] and external devices like disks, network interfaces, keyboards, and displays. Input involves receiving data from external sources, whereas output involves sending data to external destinations.

I/O in Linux:

In Linux, I/O is essential for communication between user processes, the kernel, and hardware devices. The Linux kernel provides various interfaces and mechanisms to handle I/O efficiently. Understanding the I/O subsystem is crucial for optimizing system performance and writing efficient applications.

Kernel I/O Structure:

The Linux kernel I/O structure consists of three main layers:

Block I/O Layer:

The block I/O layer is responsible for managing block devices, such as hard drives, solid-state drives (SSDs), USB drives, and other storage media. Block devices store data in fixed-size blocks, typically ranging from 512 B to 4 kB. The key components of the block I/O layer include:

Block Device Drivers: These are kernel modules that provide the interface between the kernel and specific block devices. They manage I/O requests and handle communication with the hardware. Examples include the SATA, SCSI, and NVMe drivers.

I/O Scheduler: The I/O scheduler is responsible for optimizing the order in which I/O requests are executed to minimize disk seek times and improve overall I/O performance. Popular Linux I/O schedulers include completely fair queuing (CFQ), NOOP, and Deadline.

Buffer Cache: The buffer cache is a region of the RAM used to cache data read from and write to block devices. Caching improves performance by reducing the number of actual disk accesses, especially for frequently accessed data.

Character I/O Layer:

The character I/O layer manages character devices, which transfer data character by character. Character devices include devices like keyboards, mice, terminals, and serial ports. Key components of the character I/O layer include:

Character Device Drivers: These drivers handle communication with specific character devices. For example, the TTY driver manages terminal devices (e.g., /dev/tty1).

Network Layer:

The Network Layer deals with I/O operations related to network communication. It manages data transfer in the form of packets between the computer and network devices. Key components of the network layer include:

Network Device Drivers: These drivers handle communication with network interfaces, such as Ethernet cards and Wi-Fi adapters.

Network Stack: The network stack manages the routing, flow control, and transmission of network packets. It includes components like the IP layer (IPv4 and IPv6), transport layer (TCP, UDP), and socket interface.

sysfs and procfs: Linux provides special filesystems, that is, sysfs (system file system) and procfs (process file system), to expose kernel information and configuration to user space. These filesystems enable users and system administrators to interact with kernel I/O settings and parameters using standard file I/O operations.

Kernel APIs:

The Linux kernel provides a set of application programming interfaces (APIs) that allow user-space processes to interact with the kernel and perform I/O operations. Some commonly used kernel I/O APIs include read(), write(), open(), ioctl(), and mmap().

Examples of I/O in Linux:

Reading/Writing a File:

```
#include <stdio.h>

int main() {
    FILE *file = fopen("example.txt", "r"); // Open file in read mode
    if (file != NULL) {
        char buffer[100];
        while (fgets(buffer, sizeof(buffer), file) != NULL) {
            // Process data in the buffer
            printf("%s", buffer);
        }
        fclose(file);
    }
    return 0;
}
```

In this example, the program reads a file named "example.txt" line by line using the fgets() function and prints its content on the screen.

Using I/O Redirection:

In Linux, you can redirect standard input, output, and error streams to/from files using the following symbols:

- >: redirects standard output to a file.
- <: redirects standard input from a file.
- 2>: redirects standard error to a file.

Example:

```
$ echo "Hello, I/O" > output.txt # Redirects output to a file
$ cat < input.txt                # Redirects input from a file
$ ls non_existent_file 2> error.txt # Redirects error to a file
```

Interacting with Devices:

Linux treats devices as special files in the `/dev/` directory.

For instance:

`/dev/sda`: represents the first hard disk.

`/dev/keyboard`: represents the keyboard.

Commands Related to Kernel I/O and I/O Devices:

`lsblk`: lists block devices (displays information about block devices and their partitions).

`lspci`: lists peripheral component interconnect (PCI) devices (displays information about devices connected via the PCI bus).

`lsusb`: lists USB devices (displays information about devices connected via USB).

`lsmod`: lists loaded kernel modules (displays a list of currently loaded kernel modules).

`modprobe`: loads kernel modules (loads a kernel module into the running kernel).

`rmmod`: removes kernel modules (unloads a kernel module from the running kernel).

`dmesg`: displays kernel messages (displays the kernel ring buffer and boot messages).

I/O Devices: Block and Character Devices

Introduction to I/O Devices:

I/O devices in a computer system are external peripherals that facilitate data exchange between the CPU and the outside world. Linux classifies I/O devices into two main categories: block devices and character devices. Each type has unique characteristics and usage patterns in the Linux environment.

Block Devices:

Block devices are I/O devices that transfer data in fixed-size blocks, typically ranging from 512 B to 4 kB. They include storage devices like hard disk drives, SSDs, USB drives, and redundant array of independent disks (RAID). Block devices are treated as a sequence of blocks, and data is read or written in block-sized chunks.

Characteristics of Block Devices:

Data is accessed in fixed-size blocks.

Random access is possible (reading/writing any block at any time).

Suitable for file systems and applications that require block-level access.

They have a specific capacity for storing data.

Examples: `/dev/sda` (hard disk) and `/dev/nvme0n1` (NVMe SSD).

Examples of Block Device Usage:

Mounting File Systems: When you mount a file system in Linux, it interacts with the underlying block device where the data is stored.

Example:

```
$ mount /dev/sda1 /mnt/mydrive
```

Partitioning: Partitioning involves dividing a block device into multiple logical sections to organize data efficiently.

Example (using fdisk):

```
$ sudo fdisk /dev/sda
```

Character Devices:

Character devices are I/O devices that transfer data character by character or byte by byte. They include devices such as keyboards, mice, serial ports, terminals, and audio devices. Unlike block devices, character devices do not have fixed-size blocks, and data is accessed as a stream of characters or bytes.

Characteristics of Character Devices:

- Data is accessed in individual characters or bytes.

- Sequential access is typical (reading/writing sequentially from the beginning).

- Suitable for real-time data transmission and interactive applications.

- They usually do not have a fixed capacity for storing data.

Examples: `/dev/tty` (terminal) and `/dev/ttyS0` (serial port).

Examples of Character Device Usage:

Reading Input from Keyboard: When you type on the keyboard, the characters are read from the character device representing the keyboard.

Example (reading from the terminal):

```
#include <stdio.h>

int main() {
    char buffer[100];
    fgets(buffer, sizeof(buffer), stdin); // Read input from stdin (character device)
    printf("You entered: %s\n", buffer);
    return 0;
}
```

Sending Data to a Serial Port: When communicating with external devices over a serial port, data is written to the character device representing the serial port.

Example (using write() system call):

```
#include <fcntl.h>
#include <unistd.h>

int main() {
    int fd = open("/dev/ttyS0", O_WRONLY); // Open the character device for //writing
    if (fd != -1) {
        write(fd, "Hello, Serial!", 14); // Write data to the serial port
        close(fd); // Close the file descriptor
    }
}
```

```
    }  
    return 0;  
}
```

Commands Related to Block Devices:

mount: mounts a file system (mounts a block device to a specific mount point).

umount: unmounts a file system (unmounts a mounted block device).

fdisk: partition table manipulator (used for creating and managing disk partitions).

Commands Related to Character Devices:

cat: concatenates and displays file content (can be used to read from character devices).

echo: outputs data to the terminal (can be used to write to character devices).

stty: changes and prints terminal line settings (can be used to configure character devices like serial ports).

Device Drivers in Linux Operating Systems

Introduction to Device Drivers:

Device drivers are software components that act as intermediaries between the hardware devices and the operating system. They enable the operating system to communicate and control various hardware peripherals, such as graphics cards, network adapters, USB devices, and storage devices. In Linux, device drivers play a critical role in ensuring the efficient and proper functioning of hardware devices.

Types of Device Drivers in Linux:

Linux supports two main types of device drivers:

Kernel Space Device Drivers: These drivers are directly integrated into the Linux kernel. They have direct access to kernel resources and can interact with hardware at a lower level. Kernel space drivers are generally preferred for critical and performance-sensitive tasks.

User Space Device Drivers: These drivers run in user space outside the kernel. They interact with the hardware using system calls and kernel interfaces. User-space drivers are more flexible and can be loaded and unloaded dynamically, but they may incur some performance overhead due to the user-kernel boundary crossing.

Device drivers are essential components in Linux that enable communication between the operating system and hardware devices. They play a vital role in managing various peripherals and ensuring the smooth functioning of a Linux system. Writing device drivers requires an

understanding of the Linux kernel interfaces and the specific hardware being targeted. Properly implemented device drivers can significantly enhance the performance and stability of the overall system.

Commands Related to Device Drivers:

lsmod: lists loaded kernel modules (displays a list of currently loaded kernel modules).

modinfo: shows information about a kernel module (displays information about a specific module).

insmod: inserts a module into the kernel (loads a kernel module into the running kernel).

rmmod: removes kernel modules (unloads a kernel module from the running kernel).

ls /dev: lists device files in the /dev directory (shows various device files, including block and character devices).

I/O Queuing and Interrupt Handling in Linux

Introduction to I/O Queuing:

I/O queuing is a critical mechanism in the Linux kernel that optimizes the handling of I/O requests from multiple processes and devices. It enables efficient scheduling and execution of I/O operations, reducing latency and improving system performance.

I/O Scheduling Algorithms:

In Linux, the I/O scheduler is responsible for managing the order in which pending I/O requests are serviced. Different I/O scheduling algorithms are available, and the choice of scheduler can impact system performance depending on the workload and disk characteristics. Some common I/O scheduling algorithms in Linux include:

Completely Fair Queuing (CFQ): provides a fair distribution of I/O bandwidth among processes, suitable for desktop and interactive workloads.

Deadline: guarantees time-bound I/O requests to improve responsiveness for real-time tasks.

NOOP: simple First In, First Out (FIFO)-based scheduler with minimal overhead, best suited for SSDs and devices with low latency.

Example: Changing I/O Scheduler:

You can change the I/O scheduler for a specific block device using the iosched file in the /sys filesystem.

```
# Check the current scheduler for the device (e.g., /dev/sda)
$ cat /sys/block/sda/queue/scheduler
# Change the scheduler to CFQ
$ echo cfq > /sys/block/sda/queue/scheduler
```

```
# Change the scheduler to deadline
$ echo deadline > /sys/block/sda/queue/scheduler
# Change the scheduler to NOOP
$ echo noop > /sys/block/sda/queue/scheduler
```

Interrupt Handling in Linux:

Interrupt handling is a fundamental aspect of operating systems, including Linux. When a hardware device requires attention, it raises an interrupt signal to the CPU, indicating that it needs immediate processing. Interrupt handling ensures timely response to hardware events without consuming excessive CPU resources.

Interrupt Handlers in Linux:

In Linux, interrupt handling is managed by interrupt handlers, also known as interrupt service routines. Each hardware device typically has its corresponding interrupt handler registered with the kernel.

Example: Writing an Interrupt Handler:

Below is a simplified example of an interrupt handler for a fictional device that generates an interrupt when a button is pressed.

```
#include <linux/interrupt.h>
#define BUTTON_IRQ 42
irqreturn_t button_interrupt_handler(int irq, void *dev_id) {
    // Perform necessary actions when the button is pressed
    printk(KERN_INFO "Button Pressed!\n");
    return IRQ_HANDLED;
}
static int __init init_button_device(void) {
    int result;
    // Request the IRQ line for the button
    result = request_irq(BUTTON_IRQ, button_interrupt_handler,
        IRQF_TRIGGER_FALLING, "button_device", NULL);
    if (result) {
        printk(KERN_ERR "Failed to request IRQ %d\n", BUTTON_IRQ);
        return result;
    }
    printk(KERN_INFO "Button Device Initialized\n");
    return 0;
}
static void __exit exit_button_device(void) {
    // Free the IRQ line for the button
    free_irq(BUTTON_IRQ, NULL);
    printk(KERN_INFO "Button Device Unloaded\n");
}
```

```
module_init(init_button_device);  
module_exit(exit_button_device);  
MODULE_LICENSE("GPL");
```

Commands Related to I/O Queuing and Interrupt Handling:

cat /sys/block/<device>/queue/scheduler: checks the current I/O scheduler for a block device.
echo <scheduler> > /sys/block/<device>/queue/scheduler: changes the I/O scheduler for a block device.

cat /proc/interrupts: displays interrupt statistics (shows the number of interrupts handled by each CPU).

Please note that some of these commands may require root (administrative) privileges to execute. Always exercise caution while using commands that interact with the kernel and device drivers, as improper use may lead to system instability or data loss.

Inspecting Hard Disks and Sectors in Linux Operating Systems

Introduction:

Inspecting hard disks and sectors in a Linux operating system is crucial for understanding the storage configuration and disk health and analyzing disk-related issues. Linux provides several commands and utilities to examine hard disks, partitions, and sectors, helping users and system administrators make informed decisions regarding storage management and troubleshooting.

View Disk Information:

The following commands can be used to view general information about hard disks in Linux:

lsblk: lists block devices, their mount points, and other details like size and partitions.

fdisk -l: displays partition information for all detected disks.

parted -l: shows detailed partition information using GNU Parted.

df -h: displays the disk space usage of mounted file systems.

SMART (Self-Monitoring, Analysis, and Reporting Technology):

SMART is a monitoring system used to check the health and reliability of hard drives. The smartctl command is used to access SMART data and perform tests:

`smartctl -a /dev/sdX`: shows the SMART attributes of the specified hard disk (e.g., `/dev/sda`).

`smartctl -t long /dev/sdX`: initiates a long SMART self-test on the hard disk.

Inspect Disk Partitions:

You can examine and manipulate disk partitions using the following commands:

`fdisk /dev/sdX`: interactively manages disk partitions (replace X with the appropriate disk identifier).

`gdisk /dev/sdX`: interactive GUID partition table (GPT) manipulator.

Low-Level Disk Inspection:

`dd` is a versatile utility to read or write data at the block level. Exercise caution while using these commands as they can irreversibly overwrite data:

`dd if=/dev/sdX of=/path/to/output_file`: creates an image (backup) of the entire disk (replace X with the disk identifier).

`dd if=/dev/zero of=/dev/sdX bs=1M count=1`: writes zeros to the first 1 MB of the disk.

Reading Disk Sectors:

Linux provides access to raw disk sectors using the `dd` command:

`dd if=/dev/sdX of=/path/to/output_file bs=512 count=1 skip=N`: reads the Nth sector of the disk (replace X with the disk identifier).

Inspecting hard disks and sectors in a Linux operating system is essential for managing storage, understanding disk health, and diagnosing storage-related issues. With the provided commands and utilities, users and system administrators can gain valuable insights into the storage configuration and ensure the reliability and performance of their storage devices. Always exercise caution when performing low-level operations on disks to prevent data loss or unintended consequences.

Examples:

View Disk Information:

List block devices

\$ `lsblk`

Display partition information for all disks

\$ `sudo fdisk -l`

Show detailed partition information using GNU Parted

\$ `sudo parted -l`

Display disk space usage of mounted file systems

\$ `df -h`

SMART (Self-Monitoring, Analysis, and Reporting Technology):

```
# Show the SMART attributes of the specified hard disk (e.g., /dev/sda)
$ sudo smartctl -a /dev/sda
# Perform a long SMART self-test on the hard disk
$ sudo smartctl -t long /dev/sda
```

Inspect Disk Partitions:

```
# Interactively manage disk partitions for /dev/sda (use 'd' to delete, 'n' to create, 'p'
for primary, 'w' to save)
$ sudo fdisk /dev/sda
# Interactive GUID partition table (GPT) manipulator for /dev/sda (use 'd' to delete,
'n' to create, 'w' to save)
$ sudo gdisk /dev/sda
```

Low-Level Disk Inspection:

```
# Create an image (backup) of the entire /dev/sda disk
$ sudo dd if=/dev/sda of=/path/to/output_file.img
# Write zeros to the first 1 MB of /dev/sda (use with caution, as it overwrites data)
$ sudo dd if=/dev/zero of=/dev/sda bs=1M count=1
```

Reading Disk Sectors:

```
# Read the first sector of /dev/sda and save it to the output_file
$ sudo dd if=/dev/sda of=/path/to/output_file bs=512 count=1 skip=0
```

Please note that some of these commands require superuser privileges (sudo) to execute, as they deal with low-level disk operations. Always exercise caution while using commands that involve reading or writing data at the block level, as they can cause data loss if not used correctly. Double-check the disk identifiers (/dev/sdX) before running any commands, and ensure you have a good understanding of the implications of each command before using them.

Disk Commands

lsblk, fdisk, and df Commands

1. lsblk: List Block Devices

The lsblk command is used to display information about block devices (storage devices like hard drives, SSDs, etc.) and their attributes.

Syntax:

lsblk [options] [device...]

Examples:

List all block devices:

lsblk

List block devices with filesystem and mount point information:

lsblk -f

Show only specific columns:

lsblk -o NAME,SIZE,MOUNTPOINT

Display the tree view of block devices:

lsblk -t

2. fdisk: Manipulate Disk Partitions

The fdisk command is used to create, modify, or delete disk partitions on block devices.

Syntax:

fdisk [options] device

Examples:

Start fdisk on a device (e.g., /dev/sda):

fdisk /dev/sda

Print existing partitions:

p

Create a new partition:

n

Delete a partition:

d

Write changes and exit:

w

3. df: Disk Space Usage

The df command displays information about the disk space usage of filesystems.

Syntax:

df [options] [file...]

Examples:

Show disk space usage of all mounted filesystems:

```
df -h
```

Display usage for a specific filesystem:

```
df -h /dev/sda1
```

Show disk space usage with human-readable sizes:

```
df -h
```

Display usage in 1K blocks:

```
df -k
```

Note: Make sure to use these commands with caution, especially fdisk, as modifying disk partitions can result in data loss.

The lsblk, fdisk, and df commands are powerful tools for managing and inspecting block devices, manipulating disk partitions, and monitoring disk space usage on Linux systems. Understanding how to use these commands can help you effectively manage your storage resources. Always exercise caution and double-check commands before applying changes to your system.

hwdisk and parted Linux Commands

1. hwdisk: Hardware Information

The hwdisk command is used to gather detailed hardware information about your system.

Syntax:

```
hwdisk [options]
```

Examples:

Display complete hardware information:

```
hwdisk
```

Show specific hardware category (e.g., CPU, memory, and network):

```
hwdisk --cpu
```

Save hardware information to a file:

```
hwdisk --all > hardware_info.txt
```

2. parted: Disk Partitioning

The parted command is a powerful utility for creating, resizing, and managing disk partitions.

Syntax:

```
parted [options] device
```

Examples:

Start parted on a device (e.g., /dev/sdb):

```
parted /dev/sdb
```

Create a new partition table (e.g., GPT):

```
mklabel gpt
```

Create a new partition:

```
mkpart primary ext4 0% 50%
Resize a partition:
resizepart 1 80%
Set partition flag (e.g., bootable):
set 1 boot on
Quit parted after making changes:
quit
```

Note: Both `hwinfo` and `parted` commands require superuser privileges (`sudo`) to access hardware information and manipulate disk partitions effectively.

The `hwinfo` and `parted` commands provide crucial capabilities for understanding your system's hardware and managing disk partitions. `hwinfo` allows you to gather comprehensive hardware information, whereas `parted` empowers you to create, modify, and manage disk partitions efficiently. Proper usage of these commands can aid in system administration and troubleshooting tasks. Always use these commands with care, as they deal with low-level system components.

cfdisk, sfdisk, and smartctl Linux Commands

1. cfdisk: Console-Based Disk Partitioning

The `cfdisk` command is a user-friendly console-based tool for creating, modifying, and managing disk partitions.

Syntax:

```
cfdisk device
```

Examples:

Start `cfdisk` on a device (e.g., `/dev/sdb`):

```
cfdisk /dev/sdb
```

Create a new partition:

New -> Enter partition size -> Select partition type

Delete a partition:

Select partition -> Delete

Write changes and exit:

Write -> Yes

2. sfdisk: Scriptable Disk Partitioning

The `sfdisk` command is used for scripting disk partitioning operations and creating partitions programmatically.

Syntax:

```
sfdisk [options] device
```

Examples:

Create partitions using a script file:

```
sfdisk /dev/sdb < partition_script.txt
```

Print partition table to a file:

```
sfdisk -d /dev/sdb > partition_table.txt
```

3. smartctl: SMART Monitoring and Control

The smartctl command is used to monitor and control the Self-Monitoring, Analysis, and Reporting Technology (SMART) attributes of storage devices.

Syntax:

```
smartctl [options] device
```

Examples:

Display SMART information for a device (e.g., /dev/sda):

```
smartctl -a /dev/sda
```

Short self-test on a device:

```
smartctl -t short /dev/sda
```

Check test results:

```
smartctl -l selftest /dev/sda
```

Note: SMART attributes provide insights into the health of storage devices and can help predict potential failures.

The cfdisk, sfdisk, and smartctl commands offer various capabilities for disk partitioning and storage device management. cfdisk provides an interactive console-based interface, sfdisk is used for scripted partitioning operations, and smartctl allows you to monitor the health of storage devices. Proper usage of these commands is essential to manage storage resources effectively and ensure the reliability of your system's storage devices. Always exercise caution when performing operations that involve disk partitions and storage devices.

Command Line Tool to Check the Health of a Disk Drive on a Linux System

Introduction:

The health of a disk drive is crucial for ensuring the reliability and performance of a Linux system. The smartmontools package provides a command-line interface to interact with the SMART system built into most modern hard drives and solid-state drives. It allows you to check and monitor the health and attributes of your disk drives.

Installing smartmontools:

You can install smartmontools using your distribution's package manager. For example, on Debian/Ubuntu-based systems:

```
sudo apt-get install smartmontools
```

Using smartmontools:

To check the health of a disk drive, use the smartctl command from the smartmontools package.

Syntax:

```
smartctl [options] device
```

Examples:

Display Basic Disk Information:

```
smartctl -i /dev/sda
```

This command provides general information about the disk, including its model, serial number, and firmware version.

View Comprehensive SMART Data:

```
smartctl -a /dev/sda
```

This command displays a detailed report of SMART attributes, error logs, and self-test results for the specified device.

Run Short Self-Test:

```
smartctl -t short /dev/sda
```

Initiates a short self-test on the device. You can check the test results later.

Check Self-Test Log:

```
smartctl -l selftest /dev/sda
```

This command displays the results of the self-tests that have been performed on the device.

Long Self-Test:

```
smartctl -t long /dev/sda
```

Initiates a more thorough long self-test on the device.

Check Disk Temperature:

```
smartctl -A /dev/sda | grep Temperature
```

This command retrieves the disk temperature information from the SMART data.

Interpreting Results:

SMART attributes provide insights into the health of the disk drive. Attributes with non-zero values or increasing values might indicate potential issues. You can find more information about specific attributes in the drive's documentation or online resources.

Using smartmontools and the smartctl command, you can proactively monitor the health of your disk drives, identify potential problems early, and take appropriate actions to prevent data loss and system downtime. Regularly checking and interpreting SMART data can contribute to maintaining a stable and reliable Linux system.